



GPT-2 written in Mojo

1. Matrix Multiplication

```
fn matmul_forward(  
    out: DTypePointer[dtype],  
    inp: DTypePointer[dtype],  
    weight: DTypePointer[dtype],  
    bias: DTypePointer[dtype],  
    B: Int,  
    T: Int,  
    C: Int,  
    OC: Int,  
):  
    @parameter  
    fn _calc(b: Int):  
        for t in range(T):  
            var out_bt: DTypePointer[dtype] = out + b * T * OC + t * OC  
            var inp_bt: DTypePointer[dtype] = inp + b * T * C + t * C  
  
            for o in range(OC):  
                var val: FLOAT = 0.0  
                if bias != NULL:  
                    val = bias[o]  
                var wrow: DTypePointer[dtype] = weight + o * C  
  
                @parameter  
                fn _op[width: Int](iv: Int):  
                    var t = inp_bt.load[width=width](iv) * wrow.load[width=width](iv)  
                    val += t.reduce_add[1]()  
  
                vectorize[_op, SIMD_WIDTH, unroll_factor=UNROLL_FACTOR](size=C)  
  
                out_bt[o] = val  
  
    parallelize[_calc](B)
```



GPT-2 written in Mojo

2. GeLU Activation Function

```
fn gelu_forward(out: DTypePointer[dtype], inp: DTypePointer[dtype], N: Int):  
    var s: FLOAT = sqrt(2.0 / M_PI)  
  
    var num_vectorize = N // NUM_PARALLELIZE  
  
    @parameter  
    fn _calc(ip: Int):  
        @parameter  
        fn _op[width: Int](_iv: Int):  
            var iv = ip * num_vectorize + _iv  
  
            var x = inp.load[width=width](iv)  
            var cube = 0.044715 * pow(x, 3)  
            out.store[width=width](iv, 0.5 * x * (1.0 + tanh(s * (x + cube))))  
  
        vectorize[_op, SIMD_WIDTH, unroll_factor=UNROLL_FACTOR](num_vectorize)  
  
    parallelize[_calc](NUM_PARALLELIZE)
```

LLM.

GPT-2 written in Mojo

3. Cross Entropy Loss

```
fn crossentropy_forward(  
    losses: DTypePointer[dtype],  
    probs: DTypePointer[dtype],  
    targets: DTypePointer[dtype_int],  
    B: Int,  
    T: Int,  
    Vp: Int  
):  
    # output: losses is (B,T) of the individual losses at each position  
    # input: probs are (B,T,Vp) of the probabilities  
    # input: targets is (B,T) of integers giving the correct index in logits  
  
    @parameter  
    fn _calc(b: Int):  
        for t in range(T): # todo  
            # loss = -log(probs[target])  
            var probs_bt: DTypePointer[dtype] = probs + b * T * Vp + t * Vp  
            var ix = targets[b * T + t]  
            losses[b * T + t] = -log(probs_bt[ix])  
  
    parallelize[_calc](B)
```